

Shipment Prediction Project

Chinni Jyothi Prakash

December 21, 2024

1 Introduction

This project aims to predict shipment outcomes (whether a shipment will be delayed or delivered on time) using a dataset containing shipment details. The process involved data exploration, preprocessing, model training, evaluation, and deployment of a prediction API. The final solution leverages a Random Forest model due to its superior performance and a user-friendly GUI for interacting with the API.

2 Dataset Exploration and Cleaning

2.1 Dataset Overview

The dataset contains features such as `Origin`, `Destination`, `Vehicle_Type`, `Distance`, `Weather_Conditions`, `Traffic_Conditions`, and the target variable `Delayed`. Initial inspection revealed 597 null values in the `Vehicle_Type` column.

2.2 Handling Null Values

Since `Vehicle_Type` is a categorical variable, the null values were filled with the mode of the column to ensure minimal bias.

2.3 Feature Selection

Columns like `Shipment ID`, `Shipment Date`, `Planned Delivery Date`, and `Actual Delivery Date` were excluded as they did not show significant impact on the target variable. Key columns such as `Traffic_Conditions`, `Weather_Conditions`, and `Distance (km)` were retained, as they had clear impacts on shipment delays.

3 Exploratory Data Analysis (EDA)

3.1 Visualizations

Graphical analyses revealed insights, such as a strong correlation between high/moderate traffic and shipment delays. Visualizations helped identify patterns and relationships between features and the target variable.

3.2 Unique Value Identification

Unique values in categorical columns were noted for encoding purposes.

4 Preprocessing

4.1 Categorical Encoding

Label encoders were used to convert categorical features into numerical representations.

4.2 Feature Selection

A final list of features was prepared based on their relevance to the prediction task.

4.3 Data Splitting

The dataset was split into training (80%) and testing (20%) sets for model evaluation.

5 Model Training and Evaluation

5.1 Model Selection

Three models were trained and evaluated:

- Logistic Regression
- Decision Tree
- Random Forest

5.2 Performance Metrics

Models were evaluated using metrics such as **Accuracy**, **Precision**, **Recall**, and **F1 Score**.

5.3 Results

- Logistic Regression showed relatively lower accuracy.
- Random Forest demonstrated superior performance and was chosen as the final model due to its robustness and accuracy.

Model	Accuracy	Precision	Recall	F1 Score
Logistic Regression	0.6265	0.7029	0.8521	0.7704
Random Forest	0.8825	0.9414	0.8960	0.9181
Decision Tree	0.8620	0.9064	0.9058	0.9061

Table 1: Performance Metrics of Logistic Regression, Random Forest, and Decision Tree Models

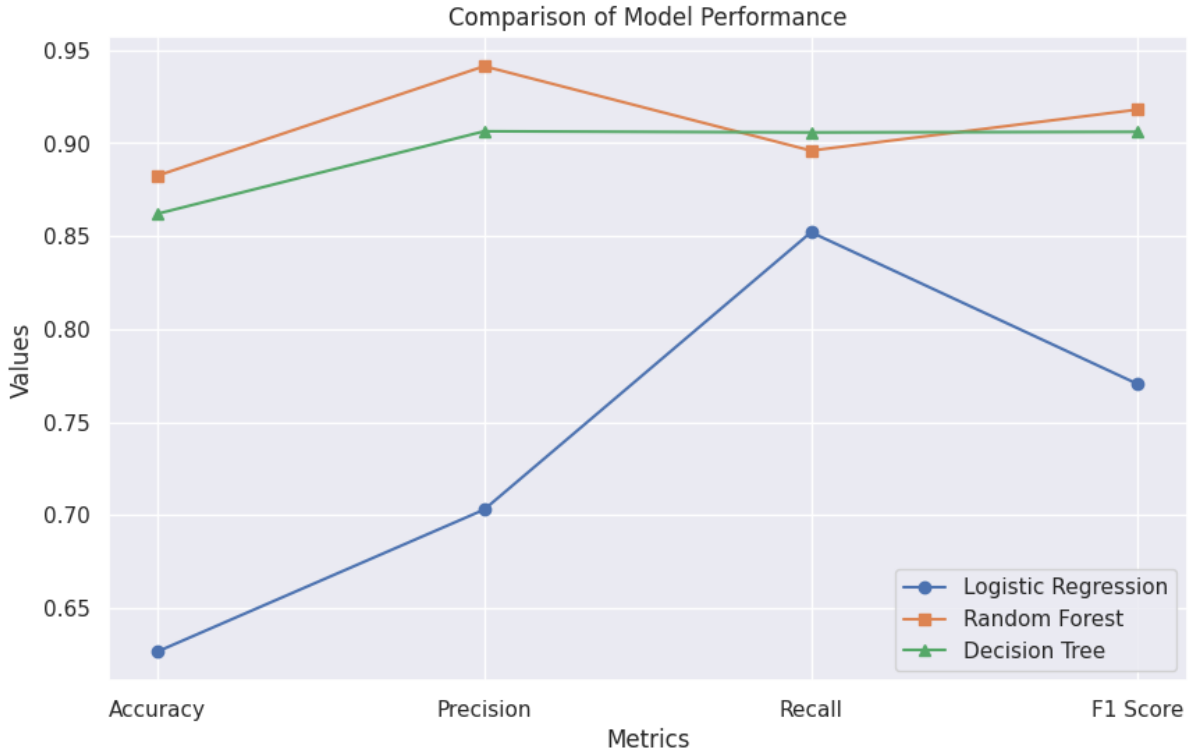


Figure 1: Performance Metrics of Logistic Regression, Random Forest, and Decision Tree Models

5.4 Visualization

Performance metrics were visualized to clearly compare the models.

5.5 Model Saving

The trained Random Forest model and label encoders were serialized using `joblib` for easy deployment in the API.

6 API Development

6.1 Setup

An environment was created in Anaconda, and libraries like `FastAPI`, `uvicorn`, and `joblib` were installed.

6.2 API Implementation

shipment_api.py was created to implement the FastAPI application. Input validation was handled using Pydantic. A /predict endpoint was developed to preprocess inputs, use the Random Forest model for predictions, and return results.

6.3 Testing the API

A client script (api_implementation.py) was created to send POST requests to the API with test data:

```
test_data = {  
    "Origin": "Jaipur",  
    "Destination": "Mumbai",  
    "Vehicle_Type": "Truck",  
    "Distance (km)": 1603.0,  
    "Weather_Conditions": "Rain",  
    "Traffic_Conditions": "Light"  
}
```

The predictions were displayed in the terminal.

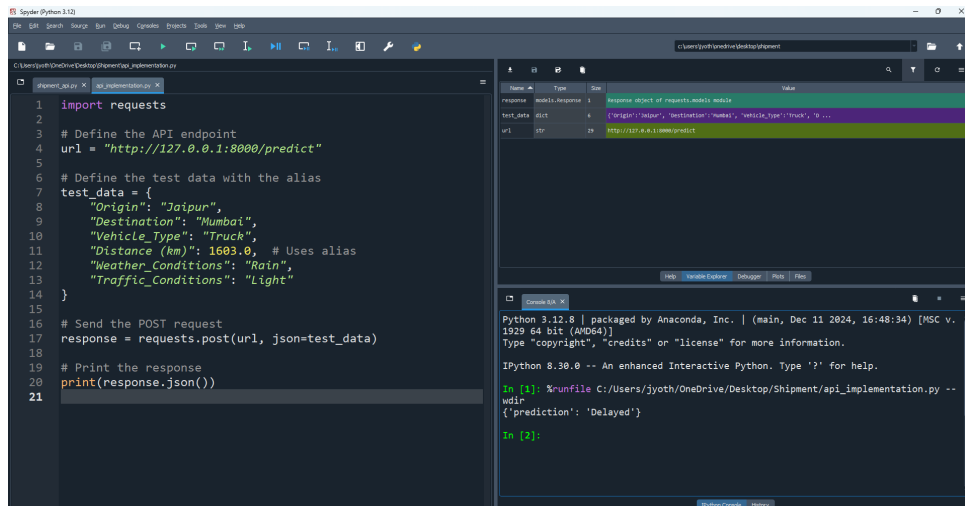


Figure 2: API Test Result

7 GUI Integration

7.1 Purpose

A small GUI was created using Tkinter to enhance usability for non-technical users.

7.2 Functionality

Users can select shipment features like Origin, Destination, Vehicle_Type, Weather_Conditions, and Traffic_Conditions from dropdown menus. The GUI interacts with the API to provide real-time predictions.

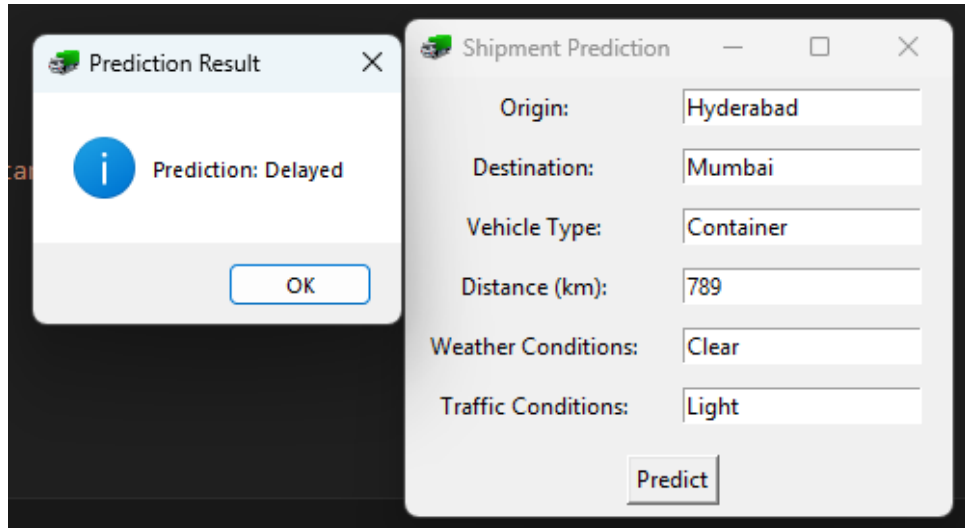


Figure 3: GUI Result

8 Key Decisions

- **Model Choice:** Random Forest was chosen due to its ability to handle diverse data types and deliver high accuracy.
- **API Framework:** FastAPI was selected for its speed, ease of use, and ability to create RESTful APIs.
- **GUI Addition:** The GUI was added to improve accessibility for users unfamiliar with APIs.

9 Conclusion

This project successfully implemented a shipment prediction system using machine learning. The Random Forest model demonstrated the best performance, and the API and GUI provide a seamless way to interact with the model. Future improvements could include hyperparameter tuning, handling more edge cases, and deploying the API to a cloud platform for broader accessibility.

```
C:\WINDOWS\system32\cmd. X + v
(ml) C:\Users\jyoth\OneDrive\Desktop\Shipment>uvicorn shipment_api:app --reload
INFO: Will watch for changes in these directories: ['C:\\Users\\jyoth\\OneDrive\\Desktop\\Shipment']
INFO: Uvicorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
INFO: Started reloader process [20592] using StatReload
INFO: Started server process [12372]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: 127.0.0.1:51003 - "POST /predict HTTP/1.1" 200 OK
WARNING: StatReload detected changes in 'interf.py'. Reloading...
INFO: 127.0.0.1:62246 - "POST /predict HTTP/1.1" 200 OK
INFO: Shutting down
INFO: Waiting for application shutdown.
INFO: Application shutdown complete.
INFO: Finished server process [12372]
INFO: Started server process [19080]
INFO: Waiting for application startup.
INFO: Application startup complete.
WARNING: StatReload detected changes in 'interf.py'. Reloading...
INFO: Shutting down
INFO: 127.0.0.1:62408 - "POST /predict HTTP/1.1" 200 OK
INFO: Waiting for application shutdown.
INFO: Application shutdown complete.
INFO: Finished server process [19080]
INFO: Started server process [19976]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: 127.0.0.1:63069 - "POST /predict HTTP/1.1" 200 OK
```

Figure 4: Terminal Result